

Extending Berti Prefetcher for Server Workloads

Phase-Aware Learning & Cross-IP Delta Correlation

Team: The Butterfly Effect

Dhruv Meena | Madhava Sriram | Saptarshi Biswas
{22b1279, 22b1233, 22b1258}@iitb.ac.in

CS683 Course Project
Indian Institute of Technology Bombay

November 2024

Outline

- 1 Introduction
- 2 Background: Berti Prefetcher
- 3 Extension 1: Cross-IP Learning
- 4 Extension 2: Phase-Aware Learning
- 5 Implementation & Configuration
- 6 Experimental Results
- 7 Milestone 2 Plans
- 8 Summary & Conclusion

Problem Statement

Challenge

Berti (MICRO'22) excels on SPEC/GAP workloads but **struggles with server applications**

Server Workload Characteristics:

- Highly irregular memory patterns
- Pointer-heavy data structures
- Frequent phase changes
- Large working sets
- Many transient IPs

Resulting Issues:

- ✗ Low coverage (sparse patterns)
- ✗ Reduced accuracy (phase pollution)
- ✗ Cache pollution
- ✗ Slow warm-up for new IPs

Project Objectives

Main Goal

Extend Berti with **two key mechanisms** to improve server workload performance

Milestone 1 (Completed):

① Cross-IP Delta Correlation

- Global delta graph
- Share patterns across IPs
- Fast warm-up

② Phase-Aware Learning

- Detect execution phases
- Separate delta histories
- Context-specific patterns

Milestone 2 (Planned):

- ③ Region-Aware Deltas
- ④ Feedback-Based Confidence

Expected Impact:

- Higher coverage
- Better accuracy
- Reduced pollution
- Faster convergence

What is Berti?

Berti: An Accurate Local-Delta Data Prefetcher (MICRO'22)

Core Concept:

- Track **delta patterns** (address differences)
- Per-IP delta history with confidence
- Issue prefetches for high-confidence deltas
- Emphasize **timeliness**: just-in-time arrival

Key Advantages:

- + More powerful than stride prefetchers
- + Hardware-friendly (low overhead)
- + Better accuracy than global methods
- + Practical for real systems

Delta Example

Address	Delta
0x1000	-
0x1040	+64
0x1080	+64
0x10C0	+64

Detected pattern: +64
⇒ Prefetch 0x1100

Berti: Key Mechanisms

1. Per-IP Delta Tracking

- Each IP maintains delta history
- Limited entries (e.g., 8 strides)
- Confidence counter per delta

2. Confidence Mechanism

- $\text{conf} += 1$ on repeated delta
- $\text{conf} -= 1$ on misprediction
- Threshold-based decision

3. Fill Level Selection

- **L1:** High confidence ($\geq 75\%$)
- **L2:** Medium confidence (35-75%)
- **Drop:** Low confidence ($< 35\%$)

Example: IP 0x400A

Delta	Conf	Level
+64	90%	L1
+128	60%	L2
-64	40%	L2
+256	20%	Drop

Timeliness:

- Track latency from issue to fill
- Adjust prefetch distance
- Avoid too-early (pollution) or too-late (miss)

Limitations of Baseline Berti

Why Berti Struggles with Server Workloads

Three fundamental limitations prevent optimal server performance:

1. Cold-Start Problem for New IPs

- New/infrequent IPs have **no history**
- Slow confidence build-up
- Server workloads have many transient IPs
- Example: Short-lived request handlers

2. Phase Changes Cause Pollution

- Single delta table per IP gets **conflicting patterns**
- Different phases overwrite each other
- Example: Database index scan → sequential scan → join operation

3. Isolated Learning

- Each IP learns **independently**

Extension 1: Cross-IP Delta Correlation

Core Idea

Build a **global delta graph** capturing stride patterns across **all IPs**

Global Delta Graph Tracks:

- ① **Confidence**: How frequently delta appears
- ② **IP Count**: Number of unique IPs using it
- ③ **Last Seen**: Recency timestamp

Ranking:

$$\text{Score} = \text{Confidence} + (\text{IP_Count} \times 5)$$

Benefits:

- + New IPs use global knowledge
- + Reduced warm-up time
- + Captures program-wide patterns
- + Handles transient IPs

Example:

- Many IPs show +64 delta
- Global graph records this
- New IP immediately benefits

Global Delta Graph Structure

Global Delta Graph

Delta	Confidence	IP Count	Last Seen	Score
+64	187	12	1000245	247
+128	143	8	1000198	183
-64	98	15	1000102	173
+256	76	5	999987	101

...

Usage: When IP has weak history, retrieve top-8 deltas from global graph

Storage: 512 entries, LRU eviction when full

Cross-IP Learning: Implementation

Adding to Global Graph:

- Sample every 4th delta (overhead control)
- Confidence $\pm= 2$ (faster build-up)
- Track contributing IPs in set
- Evict oldest if table full

Pseudocode

```
if |delta| >= MAX_STRIDE: return
if sample_count % 4 != 0: return
entry = global_graph[delta]
entry.confidence += 2
entry.ips.insert(ip)
entry.last_seen = cycle
```

Retrieving Correlations:

- Filter by confidence 3
- Score = conf + (ip_count \times 5)
- Sort by score
- Return top-8 candidates

Parameters:

- Table size: 512 entries
- Confidence threshold: 3
- Max candidates: 8
- Sample rate: 1/4

Extension 2: Phase-Aware Delta Prediction

Motivation

Programs exhibit **phase changes**: different execution contexts have different patterns

Examples of Phases:

- Database: index scan → sequential
- Web server: parse → process → respond
- ML: forward pass → backward pass
- Compiler: parse → optimize → codegen

Problem with Single Table:

- Conflicting patterns
- Old patterns reduce new confidence
- Accuracy drops

Solution: Phase Detection

- Monitor cache miss rate
- Create "phase signature"
- Detect behavior changes
- **Separate delta tables per phase**

Benefits:

- + Clean context separation
- + Higher accuracy
- + Fast re-learning on phase return



Phase Signature Structure

Phase Signature Contains:

- Miss rate percentage
- Sample count
- Last update timestamp
- Phase ID (0-7)

Detection Parameters:

- Check interval: **10K cycles**
- Min samples: 50 accesses
- Transition threshold: **8%** difference
- Max phases: 8

Matching:

- Find phase with smallest miss % diff


Example Phases

Phase	Miss %	Context
0	5	Init/Warmup
1	25	Index Scan
2	8	Sequential
3	40	Random Access

Phase-Specific Tables:

- Each (IP, Phase) has own deltas
- Tables grow independently
- Fallback to previous phase if needed


Phase-Specific vBerti Tables



Key Benefits:

- Same IP can have different patterns in different phases
- No pollution between phases
- Quick adaptation when phase returns

Complete System Architecture



Aggressive Configuration Parameters

Cross-IP Learning:

- Global delta table: **512 entries**
- Sampling rate: every **4th** delta
- Confidence increment: **+2** (vs +1)
- Confidence threshold: **3** (vs 8)
- Max candidates: **8** (vs 5)

Phase Detection:

- Check interval: **10K** cycles (vs 25K)
- Min samples: **50** (vs 100)
- Transition threshold: **8%** (vs 12%)
- Max phases: **8**

Prefetch Aggressiveness:

- L1 threshold: **60%** (vs 75%)
- L2 threshold: **25%** (vs 35%)
- L2R for low confidence (vs drop)
- More prefetch launch attempts
- Willing to use correlation

Storage Overhead:

- Main vBerti: 1024 entries
- Phase vBerti: $3 \times$ main (3072)
- Global delta: 512 entries
- **Total:** \sim **421 KB** ($4.2 \times$ baseline)

Experimental Setup

Simulator:

- ChampSim with Berti
- L1D: 32KB, 8-way
- L2: 256KB, 8-way
- LLC: 2MB, 16-way

Workloads:

- CloudSuite (server)
- SPEC2017 (single-core)
- Custom server traces


Metrics:

- IPC (Instructions Per Cycle)
- MPKI (Misses Per Kilo Instructions)
- Prefetch Accuracy & Coverage
- Memory Bandwidth


Configurations:

- Baseline Berti (original)
- Aggressive vBerti (ours)

Results: CloudSuite IPC Improvement




Results: SPEC2017 Single-Core Metrics



Key Improvements: SpeedUp +3.8%, LLC MPKI -11.6%


Performance Summary




Trade-off: Accuracy decreased by 3.8%, but net benefit is strongly positive

Analysis: Coverage vs Accuracy Trade-off

Prefetch Coverage



Prefetch Accuracy



Analysis: Coverage +10.2%, Accuracy -3.8% → More useful prefetches overall

Cross-IP Learning Effectiveness


Usage Statistics:

- 23% of prefetches used cross-IP deltas
- Highest in CloudSuite (31%)
- Critical for new/transient IPs

Example

New IP 0x5F2A:

- No local history
- Retrieved: +64, +128
- Both 95%+ accuracy
- Avoided cold-start



Insight: Small set of deltas covers majority → global learning highly effective

Planned Extension 3: Region-Aware Deltas

Motivation

Different memory regions (heap, stack, globals) have distinct access patterns

Approach:

- Partition address space
- Separate delta tables per region
- Detect boundaries via page table
- Prevent cross-region pollution

Expected Benefits:

- + Better mixed workload accuracy
- + Reduced pollution
- + Region-specific strategies

Planned Extension 4: Confidence via Feedback

Motivation

Static thresholds don't adapt to workload dynamics or system state

Approach:

- Track prefetch usefulness
- Dynamically adjust:
 - Confidence thresholds
 - Prefetch degree
 - Aggressiveness

Expected Benefits:

- + Reduced pollution
- + Better bandwidth use
- + Adaptive to pressure

References

Berti: An Accurate Local-Delta Data Prefetcher

Agustín Navarro-Torres et al., MICRO 2022

Signature Path Prefetcher

Sangkug Lym et al., MICRO 2018

Instruction Pointer Classifying Prefetcher

Samuel Pakalapati & Biswabandan Panda, ISCA 2020

ChampSim Simulator

<https://github.com/ChampSim/ChampSim>

Our Repository:

<https://github.com/Dhruv-x7x/CS683-Final-Project>

Thank You!

Questions?

Team: The Butterfly Effect

Dhruv Meena | Madhava Sriram | Saptarshi Biswas
{22b1279, 22b1233, 22b1258}@iitb.ac.in